

# HarvardX - PH125.9x CYO Project

Brad Hummel

6/23/2020

## Introduction

The objective of the PH125.9x CYO Project is to present two algorithms that predict the presence of heart disease in a patient. The objective can be broken down into three goals. The first goal is to make a prediction using the classification algorithm Decision Tree. The second goal is to make a prediction using the ensemble learning Random Forest algorithm with default parameters. The third goal is to make a prediction using Random Forest with optimized tuning parameters.

The database originated from Cleveland and is available via UCI. The database initially contained 76 attributes and a subset of 14 were made available. It consists of 13 predictor variables and 1 predicted variable called `num`. The `num` is a factor variable with 4 values and it was simplified to a binary variable where numbers greater than one are set to two (factor value 1).

Random Forest creates multiple groups, hundreds or thousands, of Decision Trees that don't use all the predictor variables in any one tree. That's one big difference between the two models, the Decision Trees use all the predictor variables. Random Forest uses hundreds of these partial Decision Trees, each tree randomly choosing a partial variable set, hence the name Random Forest.

Decision trees work well with training data but when new data is introduced it is less accurate. This is due to overfitting, a condition where using training data affects the usage of new data, the test data, ends up producing poor results. Since Random Forest uses many trees, a term called bagging, it ends up producing better results.

Creating a Random Forest consists of five steps. Step 1, a bootstrap set of the data set is created that consists of random selections of the data set by resampling. Step 2, random subsets of the bootstrap data are used to create Decision Trees. Step 3, go back to step 1 and repeat. This is done hundreds of times to create a random forest of trees. Step 4, with the forest created, evaluate each tree to predict a new data point. Step 5, the test set is used to evaluate the training set.

In this project there are a few tuning variables available for tuning the Random forest to find an optimal solution: `mtry`, `maxnodes`, `maxtrees`.

## Methods/Analysis

The data set is loaded from "<https://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/processed.cleveland.data>".

First, the data is cleaned. Column headers are added that correspond to the data set website link. The predicted attribute, `num`, is a factor variable that has four values, 1-4. To simplify the result, values 2-4 are converted to 1. Each of the 14 attributes are converted to an appropriate type. There are `?` in the data. These `?` are converted to `NA`. The attributes are identified and rows with `NA` are coerced to factor types and results in these rows being removed.

The data set is partitioned into `train` set and `test` set. The `training` set is used for training and making adjustments to the Random Forest tuning variables. The `test` set is completely ignored during this model development.

The first training was done with just defaults. Nothing was optimized. The next training was done to find the best tunable parameter `mtry`. The next training after that was done to find the best `maxnodes` with the optimized `mtry`. The final training was done to find the best `tmxdtrees` with the optimized `tmtryandmaxnodes`. Finally, these optimized paramters are used to find the optimal solution.

R version 3.6.3 is provides default packages for basic data analysis. The library `caret` is used for creating partitions. The library `dplyr` is available for any additional data wrangling. The library `ggplot2` is used for viusalization. The library `randomForest` is used for the Random Forest algorithm. A few more libraries were added in case any additional functionality is needed.

## Detailed Method/Analysis

Install packages.

```
options(warn = -1)

if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")

## Loading required package: caret
## Loading required package: lattice
## Loading required package: ggplot2
if(!require(dplyr)) install.packages("dplyr", repos = "http://cran.us.r-project.org")

## Loading required package: dplyr
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
if(!require(tidyr)) install.packages("tidyr", repos = "http://cran.us.r-project.org")

## Loading required package: tidyr
if(!require(ggplot2)) install.packages("ggplot2", repos = "http://cran.us.r-project.org")
if(!require(rattle)) install.packages("rattle", repos = "http://cran.us.r-project.org")

## Loading required package: rattle
## Loading required package: tibble
## Loading required package: bitops
## Rattle: A free graphical interface for data science with R.
## Version 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
if(!require(randomForest)) install.packages("randomForest", repos = "http://cran.us.r-project.org")

## Loading required package: randomForest
## randomForest 4.6-14
```

```

## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:rattle':
##
##     importance
## The following object is masked from 'package:dplyr':
##
##     combine
## The following object is masked from 'package:ggplot2':
##
##     margin
if(!require(RColorBrewer)) install.packages("RColorBrewer", repos = "http://cran.us.r-project.org")

## Loading required package: RColorBrewer

Load libraries.
library(caret)
library(dplyr)
library(tidyr)
library(ggplot2)
library(grid)
library(gridExtra)

##
## Attaching package: 'gridExtra'
## The following object is masked from 'package:randomForest':
##
##     combine
## The following object is masked from 'package:dplyr':
##
##     combine
library(lattice)
library(rpart)
library(rpart.plot)
library(rattle)
library(randomForest)
library(RColorBrewer)

Load the Cleveland data set.
data <- read.csv(
  "https://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/processed.cleveland.data",
  header=FALSE
)

CleanthedataSet.

Add column headings corresponding to the website's headings. Verify the columns are correct
names(data) <- c("age", "sex", "cp", "trestbps", "choi", "fbs", "restecg", "thalach", "exang", "oldpeak",
head(data)

```

```
##   age sex cp trestbps choi fbs restecg thalach exang oldpeak slope ca thai num
## 1  63  1  1    145  233  1      2    150    0    2.3    3 0.0  6.0  0
## 2  67  1  4    160  286  0      2    108    1    1.5    2 3.0  3.0  2
## 3  67  1  4    120  229  0      2    129    1    2.6    2 2.0  7.0  1
## 4  37  1  3    130  250  0      0    187    0    3.5    3 0.0  3.0  0
## 5  41  0  2    130  204  0      2    172    0    1.4    1 0.0  3.0  0
## 6  56  1  2    120  236  0      0    178    0    0.8    1 0.0  3.0  0
```

Change the  
*tnum* column's range of 1-4 to 1 and verify.

```
data$num[data$num > 1] <- 1
head(data)
```

```
##   age sex cp trestbps choi fbs restecg thalach exang oldpeak slope ca thai num
## 1  63  1  1    145  233  1      2    150    0    2.3    3 0.0  6.0  0
## 2  67  1  4    160  286  0      2    108    1    1.5    2 3.0  3.0  1
## 3  67  1  4    120  229  0      2    129    1    2.6    2 2.0  7.0  1
## 4  37  1  3    130  250  0      0    187    0    3.5    3 0.0  3.0  0
## 5  41  0  2    130  204  0      2    172    0    1.4    1 0.0  3.0  0
## 6  56  1  2    120  236  0      0    178    0    0.8    1 0.0  3.0  0
```

```
sapply(data, class)
```

```
##      age      sex      cp trestbps      choi      fbs  restecg  thalach
## "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric" "numeric"
##      exang  oldpeak      slope      ca      thai      num
## "numeric" "numeric" "numeric"  "factor"  "factor" "numeric"
```

Coerce data columns into the appropriate types and verify.

```
data <- transform(
  data,
  age=as.integer(age),
  sex=as.factor(sex),
  cp=as.factor(cp),
  trestbps=as.integer(trestbps),
  choi=as.integer(choi),
  fbs=as.factor(fbs),
  restecg=as.factor(restecg),
  thalach=as.integer(thalach),
  exang=as.factor(exang),
  oldpeak=as.numeric(oldpeak),
  slope=as.factor(slope),
  ca=as.factor(ca),
  thai=as.factor(thai),
  num=as.factor(num)
)
```

```
sapply(data, class)
```

```
##      age      sex      cp trestbps      choi      fbs  restecg  thalach
## "integer" "factor" "factor" "integer" "integer" "factor" "factor" "integer"
##      exang  oldpeak      slope      ca      thai      num
## "factor" "numeric" "factor"  "factor"  "factor" "factor"
```

Convert ? to NA.

```
data[ data == "?" ] <- NA
colSums(is.na(data))
```

```
##      age      sex      cp trestbps      choi      fbs restecg  thalach
##      0       0       0       0       0       0       0       0
##      exang  oldpeak    slope      ca      thai      num
##      0       0       0       4       2       0
```

```
summary(data)
```

```
##      age      sex      cp      trestbps      choi      fbs
## Min.   :29.00  0: 97   1: 23   Min.    : 94.0   Min.    :126.0  0:258
## 1st Qu.:48.00  1:206  2: 50   1st Qu.:120.0  1st Qu.:211.0  1: 45
## Median :56.00          3: 86   Median :130.0  Median :241.0
## Mean   :54.44          4:144   Mean    :131.7  Mean    :246.7
## 3rd Qu.:61.00          3rd Qu.:140.0  3rd Qu.:275.0
## Max.   :77.00          Max.    :200.0  Max.    :564.0
## restecg  thalach      exang      oldpeak    slope      ca      thai
## 0:151    Min.    : 71.0  0:204   Min.    :0.00   1:142   ?   : 0   ?   : 0
## 1: 4      1st Qu.:133.5  1: 99   1st Qu.:0.00   2:140   0.0 :176   3.0 :166
## 2:148     Median :153.0          Median :0.80   3: 21   1.0 : 65   6.0 : 18
##          Mean    :149.6          Mean    :1.04   2.0 : 38   7.0 :117
##          3rd Qu.:166.0          3rd Qu.:1.60   3.0 : 20  NA's: 2
##          Max.    :202.0          Max.    :6.20   NA's: 4
## num
## 0:164
## 1:139
##
##
##
##
```

```
data$thai[which(is.na(data$thai))] <- as.factor("3.0")
data <- data[!(data$ca %in% c(NA)),]
colSums(is.na(data))
```

```
##      age      sex      cp trestbps      choi      fbs restecg  thalach
##      0       0       0       0       0       0       0       0
##      exang  oldpeak    slope      ca      thai      num
##      0       0       0       0       0       0
```

```
summary(data)
```

```
##      age      sex      cp      trestbps      choi      fbs
## Min.   :29.00  0: 97   1: 23   Min.    : 94.0   Min.    :126.0  0:255
## 1st Qu.:48.00  1:202  2: 49   1st Qu.:120.0  1st Qu.:211.0  1: 44
## Median :56.00          3: 84   Median :130.0  Median :242.0
## Mean   :54.53          4:143   Mean    :131.7  Mean    :247.1
## 3rd Qu.:61.00          3rd Qu.:140.0  3rd Qu.:275.5
## Max.   :77.00          Max.    :200.0  Max.    :564.0
## restecg  thalach      exang      oldpeak    slope      ca      thai
## 0:148    Min.    : 71.0  0:201   Min.    :0.000  1:140   ?   : 0   ?   : 0
## 1: 4      1st Qu.:133.0  1: 98   1st Qu.:0.000  2:138   0.0 :176   3.0 :166
## 2:147     Median :153.0          Median :0.800  3: 21   1.0 : 65   6.0 : 18
##          Mean    :149.5          Mean    :1.052   2.0 : 38   7.0 :115
##          3rd Qu.:165.5          3rd Qu.:1.600   3.0 : 20
```

```
##           Max.      :202.0           Max.      :6.200
## num
## 0:161
## 1:138
##
##
##
##
```

Remove NA rows by coercing to factor.

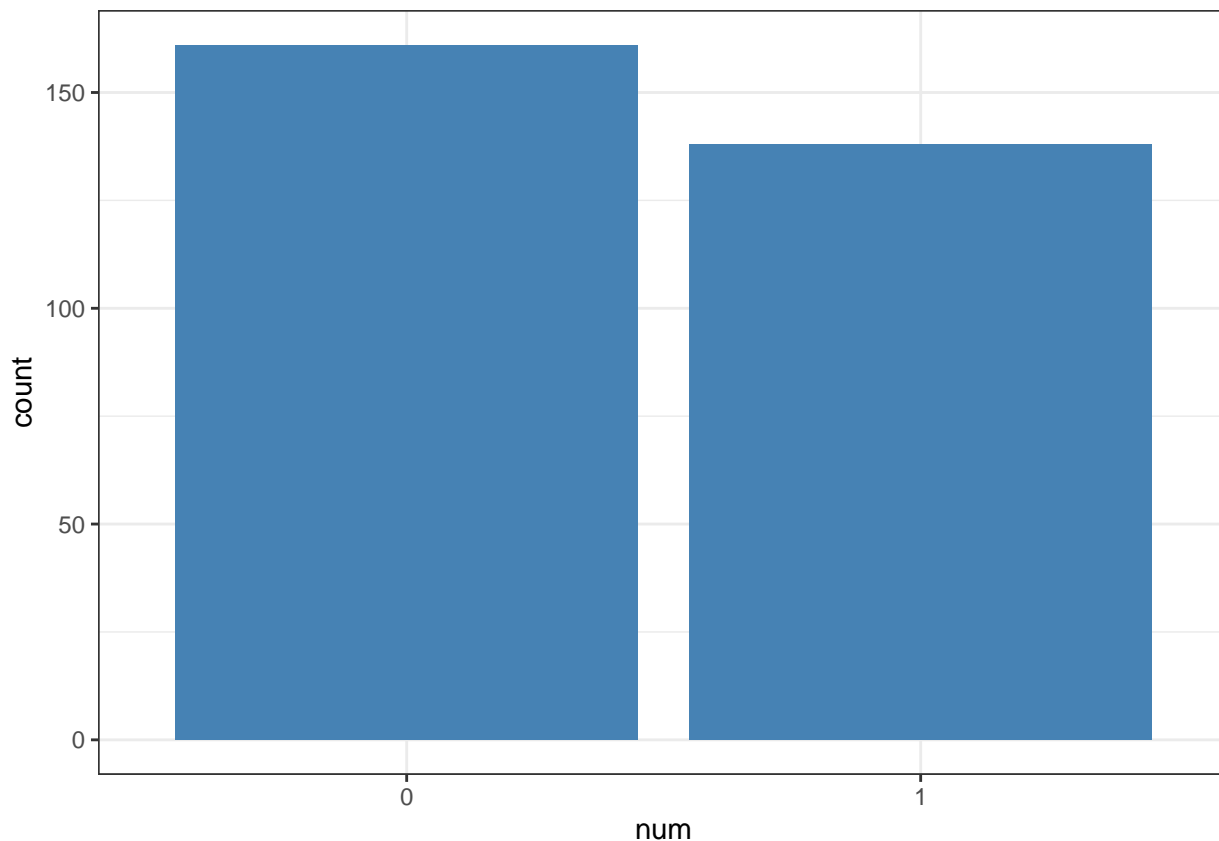
```
data$ca <- factor(data$ca)
data$thai <- factor(data$thai)
summary(data)
```

```
##      age      sex      cp      trestbps      choi      fbs
## Min.   :29.00  0: 97   1: 23   Min.    : 94.0   Min.    :126.0   0:255
## 1st Qu.:48.00  1:202  2: 49   1st Qu.:120.0   1st Qu.:211.0   1: 44
## Median :56.00           3: 84   Median :130.0   Median :242.0
## Mean   :54.53           4:143   Mean    :131.7   Mean    :247.1
## 3rd Qu.:61.00           3rd Qu.:140.0   3rd Qu.:275.5
## Max.   :77.00           Max.    :200.0   Max.    :564.0
## restecg  thalach  exang  oldpeak  slope    ca      thai
## 0:148    Min.    : 71.0  0:201   Min.    :0.000  1:140   0.0:176  3.0:166
## 1: 4     1st Qu.:133.0  1: 98   1st Qu.:0.000  2:138   1.0: 65  6.0: 18
## 2:147    Median :153.0           Median :0.800  3: 21   2.0: 38  7.0:115
##          Mean    :149.5           Mean    :1.052           3.0: 20
##          3rd Qu.:165.5           3rd Qu.:1.600
##          Max.    :202.0           Max.    :6.200
## num
## 0:161
## 1:138
##
##
##
##
```

\$Data Exploration \$

Plot proportion of disease present vs disease not present

```
data %>%
  ggplot(aes(x = num)) +
  geom_histogram(stat = 'count', fill = "steelblue") +
  theme_bw()
```



Determine number of disease vs not disease and portions

```
data_nrows <- nrow(data)
data_nrows
```

```
## [1] 299
```

```
data_present <- sum(as.numeric(data$num) == 1)
data_present
```

```
## [1] 138
```

```
data_not_present <- data_nrows - data_present
data_not_present
```

```
## [1] 161
```

```
data_present_percent <- data_present / data_nrows
data_present_percent
```

```
## [1] 0.4615385
```

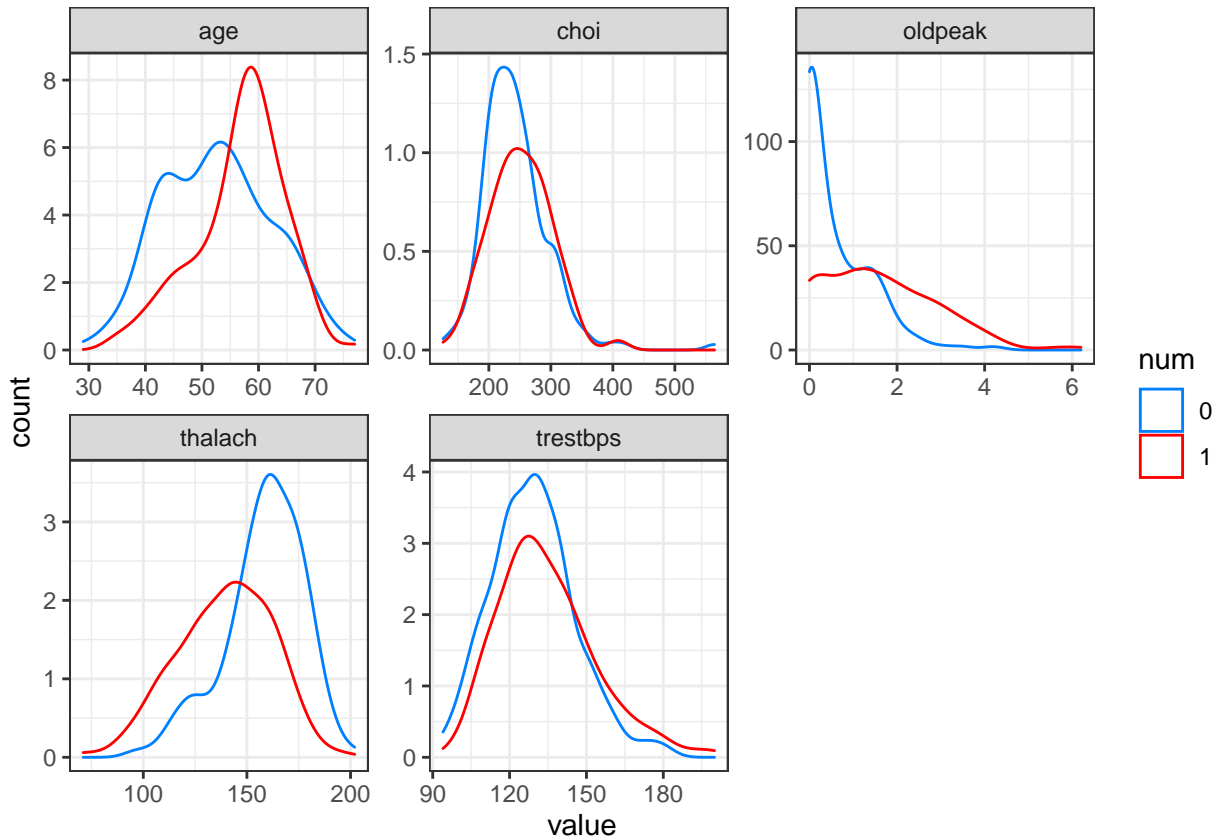
```
data_not_present_percent <- 1 - data_present_percent
data_not_present_percent
```

```
## [1] 0.5384615
```

The dimension of the database is 299 X 14. The number of patients with heart disease is 138, 46.15% of the database. The number of patients without heart disease is 161, 53.85% of the database.

Plot a few predictors.

```
data %>%
  gather(-sex, -cp, -fbs, -restecg, -exang, -slope, -ca, -thai, -num, key = "var", value = "value") %>%
  ggplot(aes(x = value, y = ..count.., colour = num)) +
  scale_color_manual(values=c("#0080FF", "#FF0000"))+
  geom_density() +
  facet_wrap(~var, scales = "free", nrow = 2) +
  theme_bw()
```



Analysis shows that as patients approach age 55 the presence of heart disease starts to exceed no presence. Cholesterol (chol) doesn't appear to have as much of an effect as advertised. The number of test subjects with heart disease is either lower or equal to test subjects without heart disease.

Create train set and test set by partitioning the data set. The train\_set and test\_set are used by the Random Forest algorithm. The dt\_train\_set and dt\_test\_set are used by the Decision Tree algorithm.

```
set.seed(123, sample.kind="Rounding")

test_index <- createDataPartition(y = data$num, times = 1, p = 0.25, list = FALSE)
train_set <- data[-test_index,]
test_set <- data[test_index,]

dt_train_set <- data[-test_index,]
dt_test_set <- data[test_index,]
```

\$Decision Tree \$ First the Descision Tree algorithm is evaluated.

Create a full Design Tree.



```
dt_full <- rpart(num ~ ., data = dt_train_set, method = "class", cp = 0)
```

The complexity table is used to control the size of the tree and find the optimal tree size. Creating a full tree with a cp setting of zero means there are no restrictions and ends up creating a complex tree.

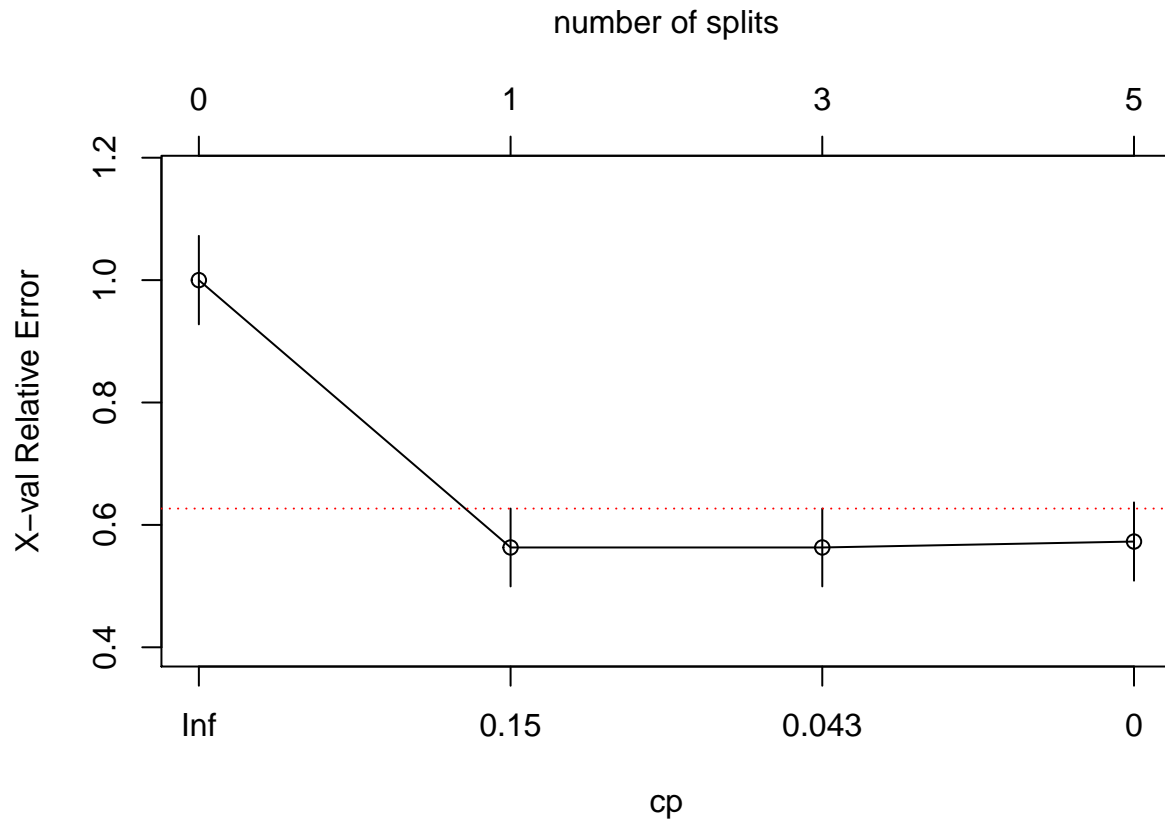
Print the complexity parameter (cp) table to help select the decision tree that minimizes misclassification error.

```
printcp(dt_full)
```

```
##
## Classification tree:
## rpart(formula = num ~ ., data = dt_train_set, method = "class",
##       cp = 0)
##
## Variables actually used in tree construction:
## [1] ca      cp      thai    thalach
##
## Root node error: 103/223 = 0.46188
##
## n= 223
##
##      CP nsplit rel error  xerror    xstd
## 1 0.485437     0   1.00000 1.00000 0.072280
## 2 0.048544     1   0.51456 0.56311 0.063601
## 3 0.038835     3   0.41748 0.56311 0.063601
## 4 0.000000     5   0.33981 0.57282 0.063953
```

Plot the cp.

```
plotcp(dt_full, lty = 3, col = 2, upper = "splits")
```



Find the best cp based on the lowest xerror.

```
bestcp <- dt_full$cptable[which.min(dt_full$cptable[, "xerror"]), "CP"]
bestcp
```

```
## [1] 0.04854369
```

Prune the bestcp.

```
dt_pruned <- prune(dt_full, cp = bestcp)
summary(dt_pruned)
```

```
## Call:
```

```
## rpart(formula = num ~ ., data = dt_train_set, method = "class",
##       cp = 0)
## n= 223
```

```
##          CP nsplit rel error   xerror   xstd
## 1 0.48543689      0 1.0000000 1.0000000 0.07228024
## 2 0.04854369      1 0.5145631 0.5631068 0.06360135
```

```
## Variable importance
```

```
##   thai thalach oldpeak   cp   sex   slope
##    45    13    12    11    10    10
```

```
## Node number 1: 223 observations,   complexity param=0.4854369
##   predicted class=0   expected loss=0.4618834   P(node) =1
##   class counts:    120   103
##   probabilities: 0.538 0.462
##   left son=2 (123 obs) right son=3 (100 obs)
```

```
## Primary splits:
##   thai   splits as LRR,      improve=30.09999, (0 missing)
##   cp     splits as LLLR,     improve=28.27550, (0 missing)
##   ca     splits as LRRR,     improve=24.09630, (0 missing)
##   thalach < 146.5 to the right, improve=21.38058, (0 missing)
##   oldpeak < 1.7  to the left, improve=19.84396, (0 missing)
## Surrogate splits:
##   thalach < 150.5 to the right, agree=0.677, adj=0.28, (0 split)
##   oldpeak < 0.75  to the left, agree=0.673, adj=0.27, (0 split)
##   cp     splits as LLLR,     agree=0.659, adj=0.24, (0 split)
##   sex     splits as LR,      agree=0.650, adj=0.22, (0 split)
##   slope   splits as LRR,     agree=0.650, adj=0.22, (0 split)
##
## Node number 2: 123 observations
##   predicted class=0 expected loss=0.2276423 P(node) =0.5515695
##   class counts:    95    28
##   probabilities: 0.772 0.228
##
## Node number 3: 100 observations
##   predicted class=1 expected loss=0.25 P(node) =0.4484305
##   class counts:    25    75
##   probabilities: 0.250 0.750
```

Predict based on the pruned decision tree.

```
dt_predict <- predict(dt_pruned, dt_test_set, type = "class")
```

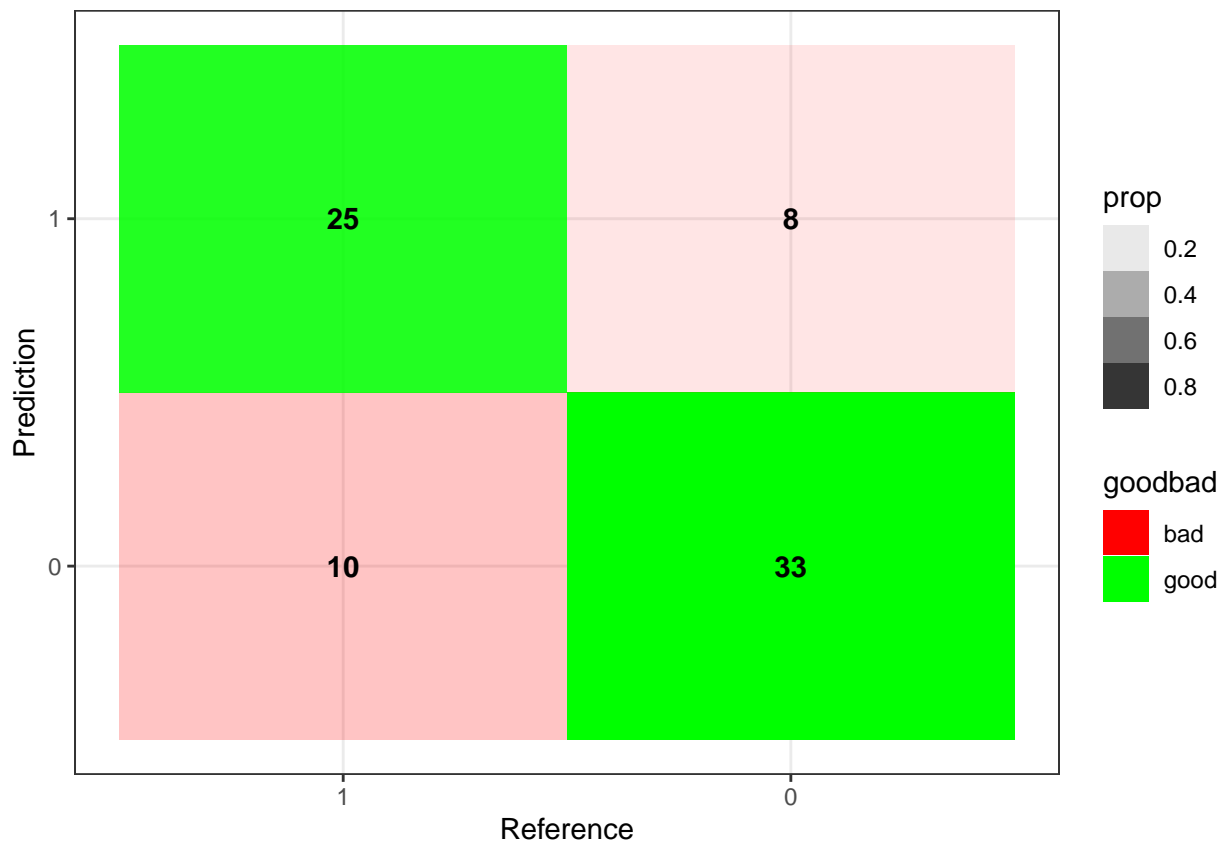
Calculate confusion matrix and plot it.

```
# Calculate confusion matrix and plot it
```

```
table <- data.frame(confusionMatrix(dt_predict, test_set$num)$table)
```

```
plotTable <- table %>%
  mutate(goodbad = ifelse(table$Prediction == table$Reference, "good", "bad")) %>%
  group_by(Reference) %>%
  mutate(prop = Freq/sum(Freq))
```

```
ggplot(data = plotTable, mapping = aes(x = Reference, y = Prediction, fill = goodbad, alpha = prop)) +
  geom_tile() +
  geom_text(aes(label = Freq), vjust = .5, fontface = "bold", alpha = 1) +
  scale_fill_manual(values = c(good = "green", bad = "red")) +
  theme_bw() +
  xlim(rev(levels(table$Reference)))
```



Calculate the classification error, subtract from 1 to get accuracy.

```
dt_accuracy <- 1 - round((table$Freq[2] + table$Freq[3]) / (table$Freq[1] + table$Freq[2] + table$Freq[3]))
dt_accuracy
```

```
## [1] 0.763
```

\$Random Forest Algorithm \$ Now let's see how the Random Forest algorithm performs.

K-fold cross validation is handled by the the trControl function. cv is the tyoe of method used for resampling the dataset. number is the number of folders. The type of search is a grid of variables to try.

```
set.seed(123, sample.kind="Rounding")
```

```
trControl <- trainControl(method = "cv",
                          number = 10,
                          search = "grid")
```

Train the default Random Forest model.

```
rf_default <- train(num ~ .,
                   data = train_set,
                   method = "rf",
                   metric = "Accuracy",
                   trControl = trControl)
```

```
rf_default
```

```
## Random Forest
```

```
##
## 223 samples
## 13 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 201, 201, 201, 201, 200, 201, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.8213439 0.6368605
## 11 0.7719368 0.5369794
## 20 0.7766798 0.5461955
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

Make prediction for the Random Forest with default parameters.

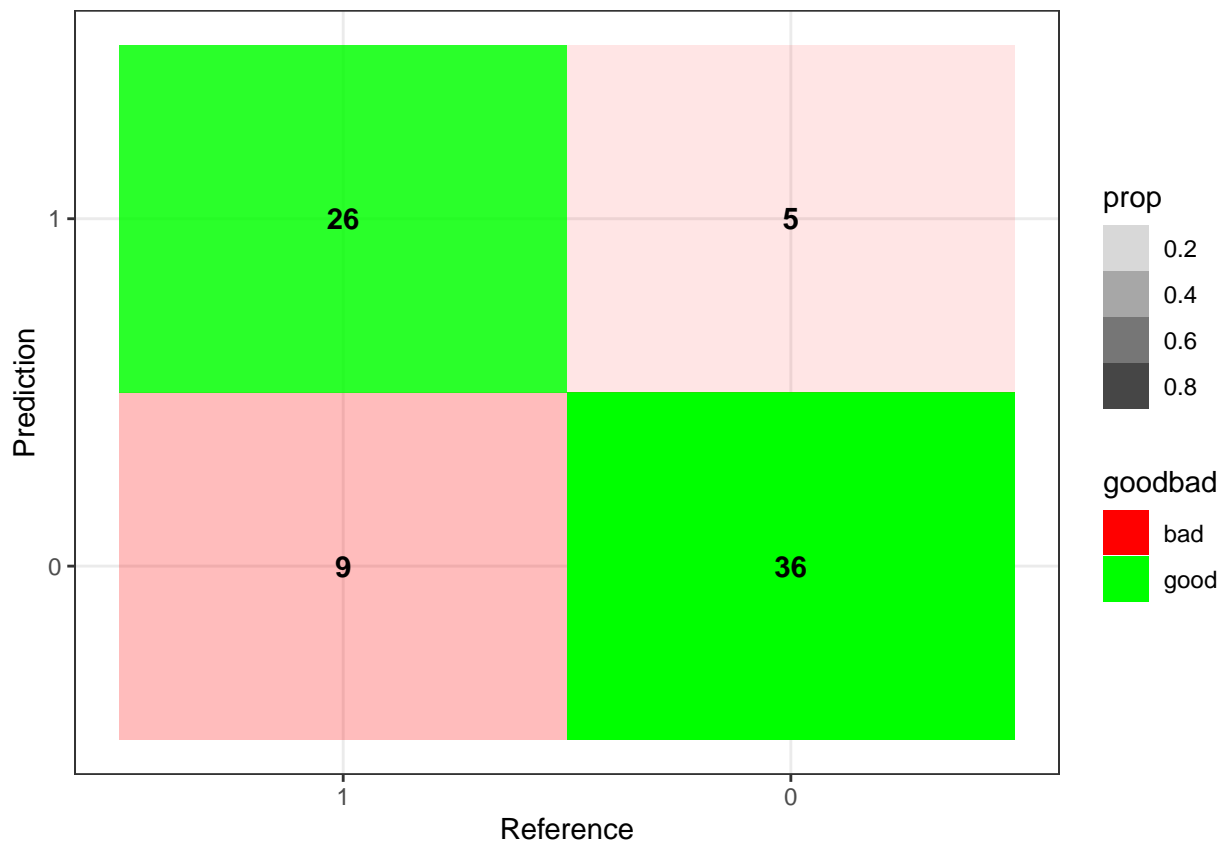
```
pred_default <- predict(rf_default, test_set)
```

Calculate the Confusion Matrix and plot it.

```
table <- data.frame(confusionMatrix(pred_default, test_set$num)$table)
```

```
plotTable <- table %>%
  mutate(goodbad = ifelse(table$Prediction == table$Reference, "good", "bad")) %>%
  group_by(Reference) %>%
  mutate(prop = Freq/sum(Freq))
```

```
ggplot(data = plotTable, mapping = aes(x = Reference, y = Prediction, fill = goodbad, alpha = prop)) +
  geom_tile() +
  geom_text(aes(label = Freq), vjust = .5, fontface = "bold", alpha = 1) +
  scale_fill_manual(values = c(good = "green", bad = "red")) +
  theme_bw() +
  xlim(rev(levels(table$Reference)))
```



Calculate classification error to get accuracy.

```
rf_accuracy <- 1 - round((table$Freq[2] + table$Freq[3]) / (table$Freq[1] + table$Freq[2] + table$Freq[3]))
rf_accuracy
```

```
## [1] 0.816
```

Calculate the variable importance and plot it.

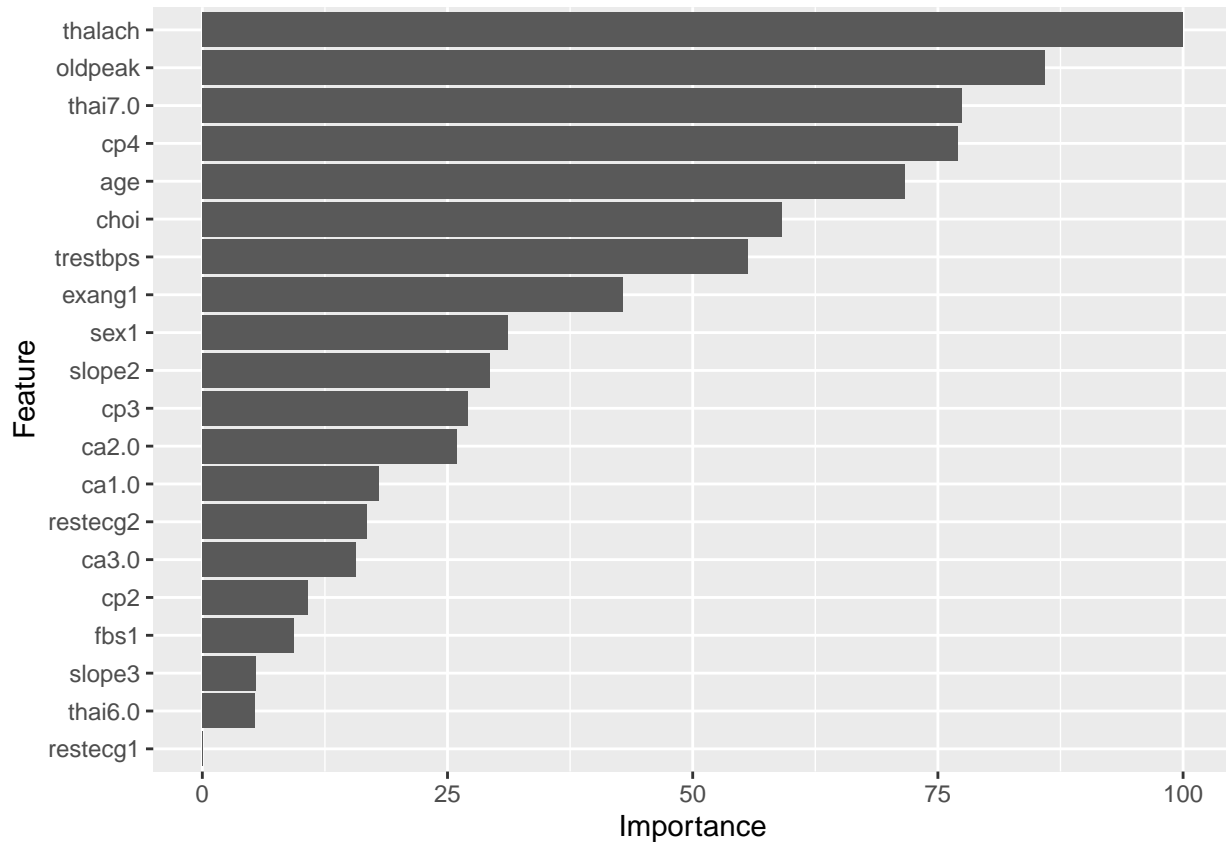
```
varImp(rf_default)
```

```
## rf variable importance
```

```
##
## Overall
## thalach 100.000
## oldpeak 85.928
## thai7.0 77.437
## cp4 76.983
## age 71.645
## choi 59.108
## trestbps 55.577
## exang1 42.905
## sex1 31.194
## slope2 29.331
## cp3 27.043
## ca2.0 25.960
## ca1.0 18.044
## restecg2 16.784
## ca3.0 15.654
```

```
## cp2      10.747
## fbs1      9.363
## slope3    5.454
## thai6.0   5.317
## restecg1  0.000
```

```
ggplot2::ggplot(varImp(rf_default))
```



Find the best mtry.

```
tuneMtry <- expand.grid(.mtry = c(1: 10))
rf_mtry <- train(num ~ .,
  data = train_set,
  method = "rf",
  metric = "Accuracy",
  tuneGrid = tuneMtry,
  trControl = trControl,
  importance = TRUE,
  nodesize = 14,
  ntree = 300)
```

```
rf_mtry
```

```
## Random Forest
##
## 223 samples
## 13 predictor
## 2 classes: '0', '1'
```

```

##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 201, 200, 201, 201, 201, 200, ...
## Resampling results across tuning parameters:
##
##   mtry  Accuracy   Kappa
##   1    0.8158103  0.6226722
##   2    0.8154150  0.6242924
##   3    0.7978261  0.5895336
##   4    0.7974308  0.5883713
##   5    0.7837945  0.5598265
##   6    0.7928854  0.5793295
##   7    0.7794466  0.5504426
##   8    0.8017787  0.5970459
##   9    0.7796443  0.5518354
##  10    0.7750988  0.5433145
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 1.

best_mtry <- rf_mtry$bestTune$mtry
best_mtry

## [1] 1

Find best maxnodes.

maxNodeList <- list()
tuneMtry <- expand.grid(.mtry = best_mtry)

for (mn in c(5: 15)) {
  set.seed(123, sample.kind="Rounding")
  rf_maxnode <- train(num ~ .,
                      data = train_set,
                      method = "rf",
                      metric = "Accuracy",
                      tuneGrid = tuneMtry,
                      trControl = trControl,
                      importance = TRUE,
                      nodesize = 14,
                      maxnodes = mn,
                      ntree = 300)

  current_iteration <- toString(mn)
  maxNodeList[[current_iteration]] <- rf_maxnode
}

results_maxnodes <- resamples(maxNodeList)
summary(results_maxnodes)

##
## Call:
## summary.resamples(object = results_maxnodes)
##
## Models: 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15
## Number of resamples: 10
##

```

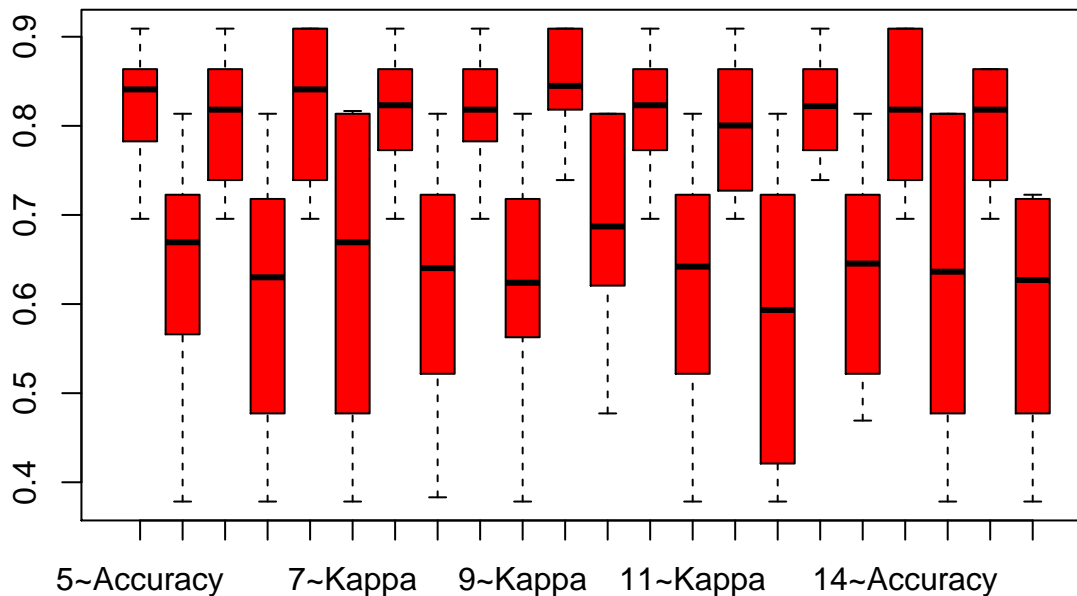


```
## Accuracy
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## 5  0.6956522 0.7915020 0.8409091 0.8217391 0.8636364 0.9090909    0
## 6  0.6956522 0.7475296 0.8181818 0.8083004 0.8636364 0.9090909    0
## 7  0.6956522 0.7475296 0.8409091 0.8264822 0.9090909 0.9090909    0
## 8  0.6956522 0.7727273 0.8231225 0.8126482 0.8636364 0.9090909    0
## 9  0.6956522 0.7826087 0.8181818 0.8169960 0.8636364 0.9090909    0
## 10 0.7391304 0.8181818 0.8448617 0.8438735 0.8977273 0.9090909    0
## 11 0.6956522 0.7727273 0.8231225 0.8083004 0.8636364 0.9090909    0
## 12 0.6956522 0.7386364 0.8003953 0.8037549 0.8636364 0.9090909    0
## 13 0.7391304 0.7840909 0.8221344 0.8213439 0.8636364 0.9090909    0
## 14 0.6956522 0.7475296 0.8181818 0.8175889 0.9090909 0.9090909    0
## 15 0.6956522 0.7475296 0.8181818 0.8037549 0.8636364 0.8636364    0
```

```
##
## Kappa
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## 5  0.3783784 0.5797007 0.6693192 0.6358496 0.7215040 0.8135593    0
## 6  0.3783784 0.4883893 0.6300170 0.6078112 0.7179487 0.8135593    0
## 7  0.3783784 0.4883893 0.6693192 0.6450646 0.8135593 0.8166667    0
## 8  0.3831418 0.5276680 0.6403432 0.6186463 0.7226891 0.8135593    0
## 9  0.3783784 0.5635627 0.6239041 0.6267822 0.7179487 0.8135593    0
## 10 0.4772727 0.6253533 0.6871021 0.6834265 0.7908418 0.8135593    0
## 11 0.3783784 0.5295157 0.6419932 0.6101668 0.7215040 0.8135593    0
## 12 0.3783784 0.4540009 0.5933637 0.5989902 0.7215040 0.8135593    0
## 13 0.4692308 0.5496377 0.6454297 0.6365331 0.7215040 0.8135593    0
## 14 0.3783784 0.4883893 0.6363388 0.6292138 0.8135593 0.8135593    0
## 15 0.3783784 0.4961659 0.6270115 0.6017078 0.7179487 0.7226891    0
```

Plot best maxnodes.

```
boxplot(results_maxnodes$values[1:10,2:23], col = "red")
```



Find the best maxtrees.

```
maxTreeList <- list()
tuneMtry <- expand.grid(.mtry = best_mtry)
```

```

for (nt in c(250, 300, 350, 400, 450, 500, 550, 600, 800, 1000, 2000)) {
  set.seed(123)
  rf_maxtrees <- train(num ~ .,
                      data = train_set,
                      method = "rf",
                      metric = "Accuracy",
                      tuneGrid = tuneMtry,
                      trControl = trControl,
                      importance = TRUE,
                      nodesize = 14,
                      maxnodes = 10,
                      ntree = nt)

  key <- toString(nt)
  maxTreeList[[key]] <- rf_maxtrees
}
results_tree <- resamples(maxTreeList)
summary(results_tree)

```

```

##
## Call:
## summary.resamples(object = results_tree)
##
## Models: 250, 300, 350, 400, 450, 500, 550, 600, 800, 1000, 2000
## Number of resamples: 10
##
## Accuracy
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## 250 0.7391304 0.7840909 0.8221344 0.8304348 0.8977273 0.9090909    0
## 300 0.7391304 0.8181818 0.8448617 0.8438735 0.8977273 0.9090909    0
## 350 0.7391304 0.7840909 0.8448617 0.8349802 0.8977273 0.9090909    0
## 400 0.7391304 0.7751976 0.8409091 0.8306324 0.8977273 0.9090909    0
## 450 0.7391304 0.7751976 0.8409091 0.8306324 0.8977273 0.9090909    0
## 500 0.7391304 0.7475296 0.8409091 0.8262846 0.8977273 0.9090909    0
## 550 0.7391304 0.7475296 0.8409091 0.8262846 0.8977273 0.9090909    0
## 600 0.7391304 0.7588933 0.8636364 0.8353755 0.8977273 0.9090909    0
## 800 0.6956522 0.7915020 0.8409091 0.8306324 0.8636364 0.9090909    0
## 1000 0.6956522 0.7915020 0.8636364 0.8351779 0.8636364 0.9090909    0
## 2000 0.6956522 0.7475296 0.8181818 0.8128458 0.8636364 0.9090909    0
##

```

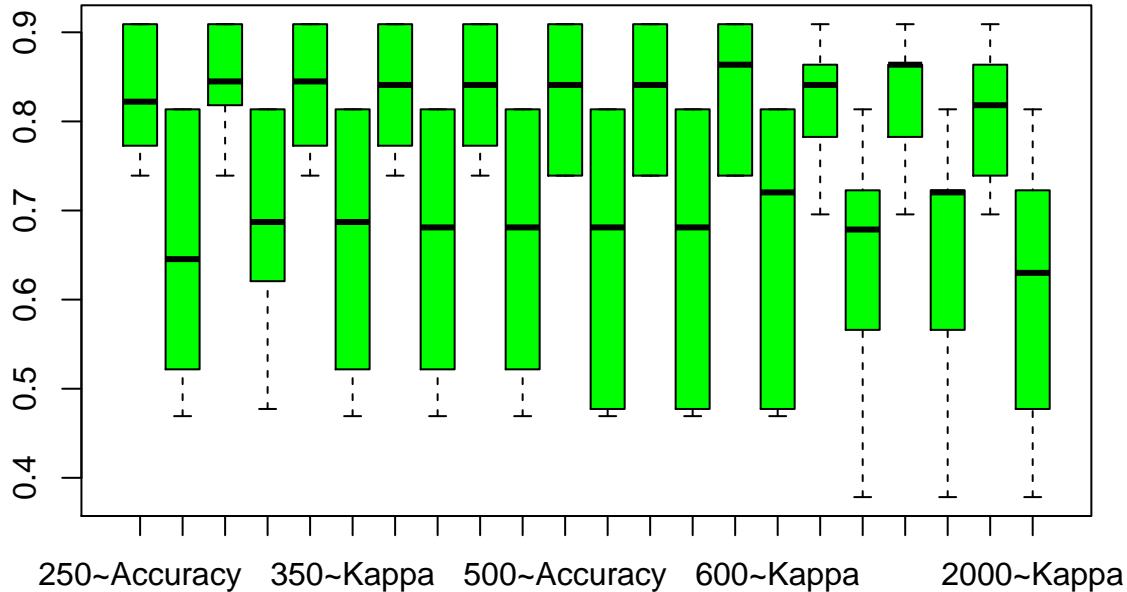
```

## Kappa
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## 250 0.4692308 0.5496377 0.6454297 0.6555802 0.7908418 0.8135593    0
## 300 0.4772727 0.6253533 0.6871021 0.6834265 0.7908418 0.8135593    0
## 350 0.4692308 0.5511404 0.6871021 0.6645158 0.7908418 0.8135593    0
## 400 0.4692308 0.5328138 0.6810167 0.6559681 0.7908418 0.8135593    0
## 450 0.4692308 0.5328138 0.6810167 0.6559681 0.7908418 0.8135593    0
## 500 0.4692308 0.4883893 0.6810167 0.6470916 0.7908418 0.8135593    0
## 550 0.4692308 0.4883893 0.6810167 0.6470916 0.7908418 0.8135593    0
## 600 0.4692308 0.5177906 0.7203189 0.6667125 0.7908418 0.8135593    0
## 800 0.3783784 0.5828616 0.6786465 0.6570277 0.7226891 0.8135593    0
## 1000 0.3783784 0.5843644 0.7203189 0.6659633 0.7226891 0.8135593    0
## 2000 0.3783784 0.4943182 0.6300170 0.6206919 0.7226891 0.8135593    0

```

Plot best maxtrees.

```
boxplot(results_tree$values[1:10,2:23], col = "green")
```



Now that we have the optimized tuning parameters, generate final fit.

```
set.seed(123)

tuneMtry <- expand.grid(.mtry = best_mtry)

rf_fit <- train(
  num ~ .,
  data = train_set,
  method = "rf",
  metric = "Accuracy",
  trControl = trControl,
  tuneGrid = tuneMtry,
  importance = TRUE,
  nodesize = 14,
  maxnodes = 10,
  ntree = 800)

rf_fit

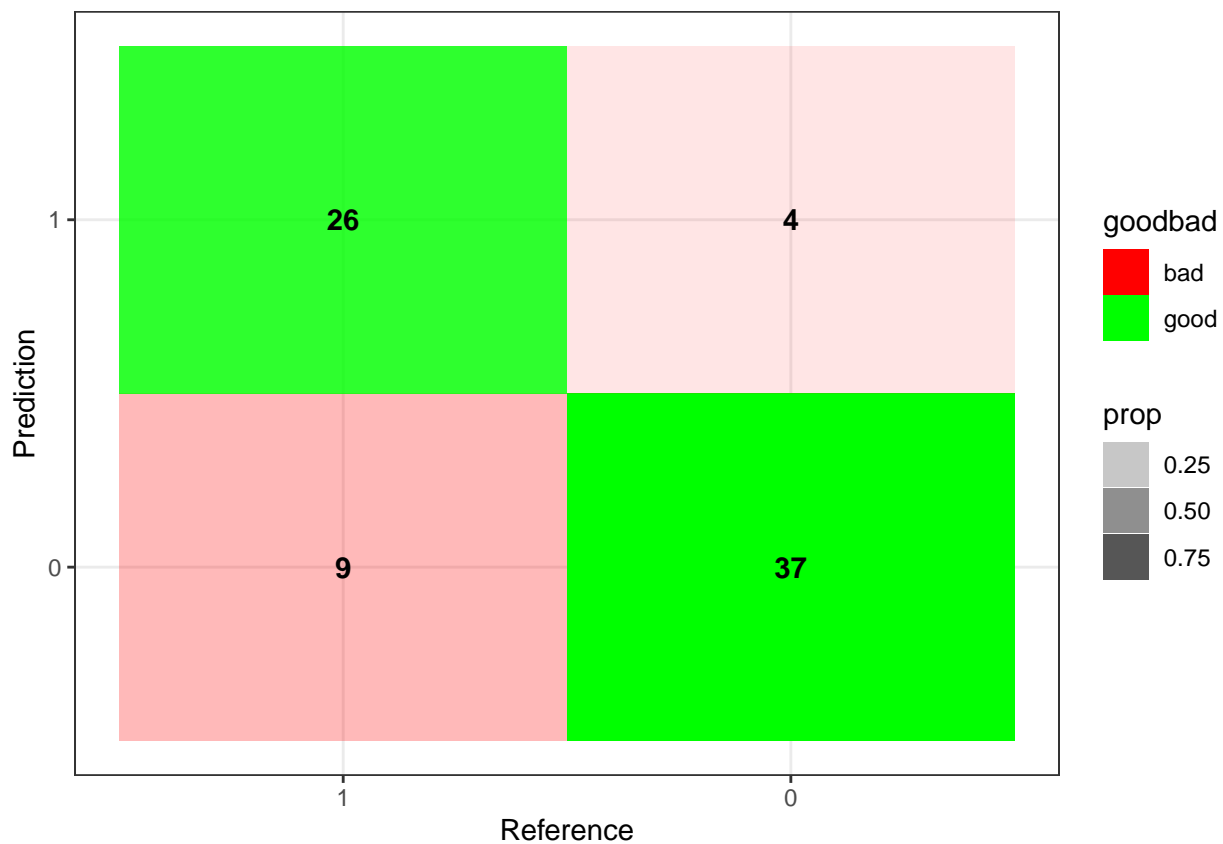
## Random Forest
##
## 223 samples
## 13 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 201, 201, 201, 201, 200, 201, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8306324 0.6570277
```

```
##
## Tuning parameter 'mtry' was held constant at a value of 1
Make a prediction with the final fit.
prediction <- predict(rf_fit, test_set)

Plot Confusion Matrix.
table <- data.frame(confusionMatrix(prediction, test_set$num)$table)

plotTable <- table %>%
  mutate(goodbad = ifelse(table$Prediction == table$Reference, "good", "bad")) %>%
  group_by(Reference) %>%
  mutate(prop = Freq/sum(Freq))

ggplot(data = plotTable, mapping = aes(x = Reference, y = Prediction, fill = goodbad, alpha = prop)) +
  geom_tile() +
  geom_text(aes(label = Freq), vjust = .5, fontface = "bold", alpha = 1) +
  scale_fill_manual(values = c(good = "green", bad = "red")) +
  theme_bw() +
  xlim(rev(levels(table$Reference)))
```



Calculate the classification error, subtract from 1 to get accuracy. Sum up the false positives and false negatives and divide by the total of the confusion matrix.

```
rf_accuracy <- 1 - round((table$Freq[2] + table$Freq[3]) / (table$Freq[1] + table$Freq[2] + table$Freq[3] + table$Freq[4]))
rf_accuracy
```

```
## [1] 0.829
```

Determine variable importance and plot.

```
varImp(rf_fit)
```

```
## rf variable importance
```

```
##
```

```
##          Importance
```

```
## thai7.0      100.000
```

```
## cp4           97.708
```

```
## oldpeak       82.482
```

```
## thalach        80.444
```

```
## exang1         76.470
```

```
## ca2.0          74.818
```

```
## slope2         73.831
```

```
## cp3            73.198
```

```
## ca3.0          71.603
```

```
## sex1           62.334
```

```
## age            55.849
```

```
## cp2            44.157
```

```
## ca1.0          44.068
```

```
## trestbps       41.496
```

```
## thai6.0        27.324
```

```
## restecg2       25.872
```

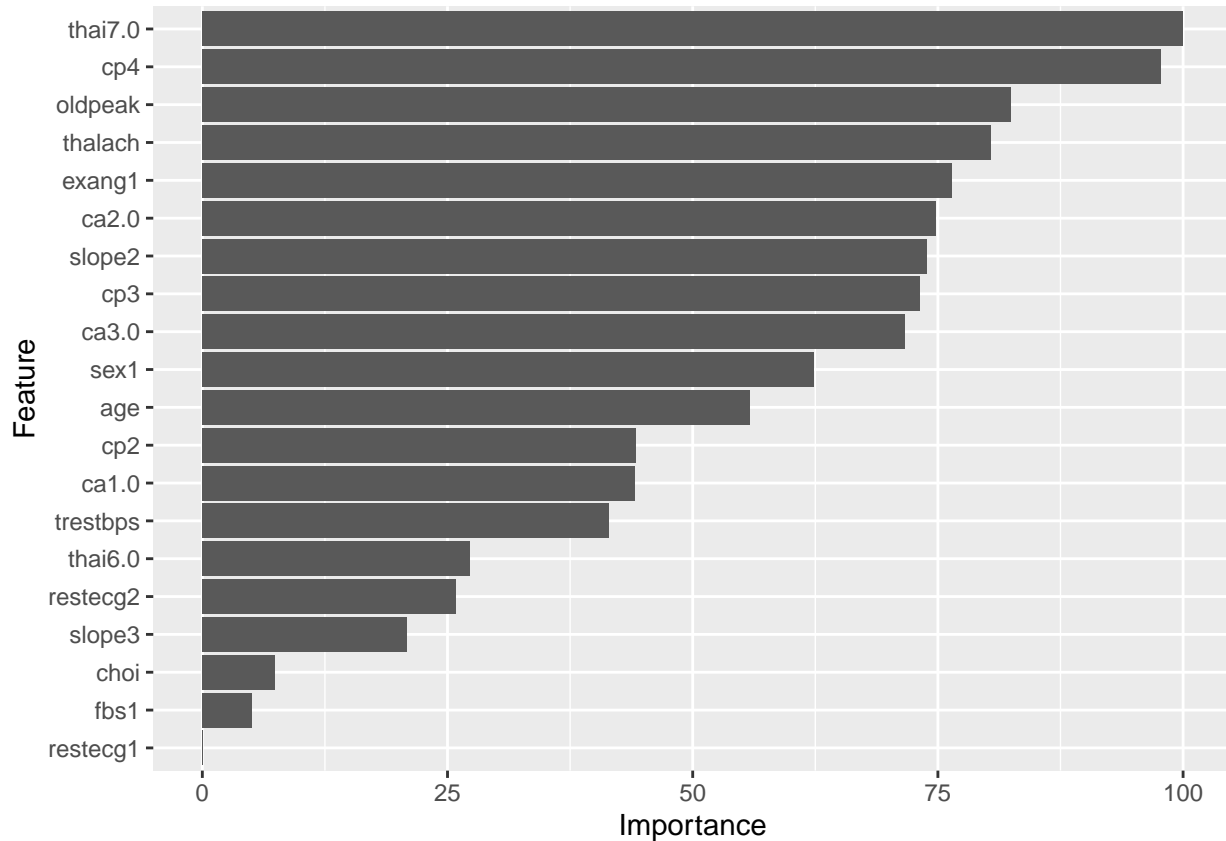
```
## slope3         20.818
```

```
## choi           7.411
```

```
## fbs1           5.033
```

```
## restecg1       0.000
```

```
ggplot2::ggplot(varImp(rf_fit))
```



## Results

The Random Forest algorithm produced better results than the Design Tree algorithm.

Accuracy: Data Tree model with default parameters 0.763; Random Forest model with default parameters 0.816; Random Forest model with optimized tuning parameters 0.829.

The complexity parameter (cp) table showed that the best cp to prune was the second one, with a cp of 0.048544 and the lowest xerror of 0.56311. Actually a third cp had the same xerror value but 2 was chosen first.

Three tuning parameters were used to generate an optimized fit for Random Forest: mtry, maxnodes and maxtrees. The optimized tuning parameters were: a) mtryof1; b) maxnodesof11; c) maxtreesof350.

The tuning parameter mtry is the number of randomly chosen variables at each split in the tree. The tuning parameter maxnodes is the number of terminal nodes. The tuning parameter maxtree is the number of trees to build.

The mtry tuning variable of 1 was chosen out the range 1 to 10.

The default variable is normally

$$\sqrt{\text{number} - \text{of} - \text{variables}} \quad (1)$$

Lower mtry values tend to lead to less correlated trees which increases stability. Moderate values have more of an effect when mtry values are low. However, lower mtry values can perform worse since less optimal values are used. Higher values were attempted but the accuracy was always worse.

The optimal maxtrees found was 350, a little lower than expected. The maxtrees don't appear as tunable but from the literature it is recommended they are high. Values from

500 to 2000 were attempted as well as low as 50 and didn't appear to have very much effect. Intuitively one would think that a higher number of trees would lead to more stability and lead to better results.

The optimal maxnodes found was 11. This is the maximum number of terminal nodes. Changes to this parameter did not seem to have much effect.

The nodesize that was used was 14, the number of variables. The typical default is 1 for classification but performance improved with 14.

The variable importance shows the importance of variables in the prediction rule. The 4 most important variables for the Random Forest default approach are thaalach thai7.0 (100%), oldpeak (86%), thai7.0 (77%), and cp4 (77%). The least important variables for the Random Forest default approach are restecgi (0%), thai6.0 (5%), slope (5%), and fbsi (9%).

The variable importance changed from the the Random Forest default approach to the optimized approach. The 4 most important variables for the Random Forest default approach are thai7.0 (100%), cp4 (98%), oldpeak (82%), and thaalach (80%). The least important variables for the Random Forest default approach are restecgi (0%), fbsi (5%), chol (7%), and slope3 (20%).

The higher number of trees the more stable the variable importance. Increasing mtry will increase the magnitude of variable of importance and affects splitting.

The database was partitioned into 75% training and 25% test. A few other partitions were tried: 0.2, 0.3, 0.4. The 0.25 partition index produced the most accurate results out of the partitions attempted.

## Conclusion

Decision Tree is a weak learner and suffers from overfitting. Overfitting means that it can get overtrained with the training set and when finally evaluated with the test set, it doesn't perform very well.

The Random Forest uses many decision trees (hundreds or thousands) to create an ensemble of trees and doesn't overfit too often. This leads to a generalization of a group of weak learning decision trees to create a stronger learning algorithm. This has been shown by the accuracy results above.

The train function was used for training Random Forests. The randomForest function based on Breiman's algorithm produces better results and would be interesting future work.

Ensemble learning algorithms like Gradient Boosting Model (GBM) and Extreme Gradient Boosting (XGB) model would also be good candidates for future evaluation. GBM caused R Studio to crash and would like to look into why it was occurring. XGB produced good results and more time is needed analyze it further.